






—●— oracle (ceiling) —●— mixture, $w = 1/K$
 —●— mixture, $w = G$ —●— expected (sample one)



per-token NLL is not comparable across tokenizers; the difference from each model's own pooled baseline is