

base model at $\sigma/\hat{\sigma}_0^* = 0.50$, round 11, seed 0

Qwen2.5-1.5B	53.5%	70.5%	2.0152	1.9459	1.8857
Qwen2.5-3B	51.9%	70.3%	1.9438	1.8809	1.8226
Qwen2.5-7B	50.1%	69.8%	1.8702	1.8144	1.7588
Llama-3.2-1B	45.0%	71.1%	1.4826	1.4446	1.3982
Gemma-3-1B †	86.1%	68.9%	1.7471	1.4465	1.3981
	captured	agreement	pooled	router	oracle

Shaded columns are ratios and compare across families. Per-token NLL does not: Qwen ~152k tokens, Llama 128k, Gemma 262k.

† the Gemma generalist had not converged at twelve rounds, so its margin is inflated.