

Higher is better in every column (↑). Bold marks the best setting in a column.

		router – pooled	oracle – pooled	captured	argmax agreement
competitive reach $\sigma/\hat{\sigma}_0^*$	0.70 $n=4$	+0.0407 ±0.0034	+0.1269 ±0.0041	32.1% ±1.7	65.9% ±3.3
	0.60 $n=4$	+0.0673 ±0.0040	+0.1351 ±0.0028	49.8% ±1.9	71.8% ±2.2
	0.55 $n=4$	+0.0714 ±0.0051	+0.1337 ±0.0053	53.4% ±1.8	71.4% ±1.2
	0.50 $n=4$	+0.0787 ±0.0079	+0.1348 ±0.0037	58.3% ±4.4	72.3% ±2.3
	0.45 $n=4$	+0.0744 ±0.0091	+0.1274 ±0.0080	58.2% ±4.0	70.0% ±3.7
	0.40 $n=4$	+0.0798 ±0.0113	+0.1295 ±0.0079	61.4% ±5.3	72.8% ±3.7
	0.30 $n=4$	+0.0808 ±0.0073	+0.1293 ±0.0029	62.4% ±4.7	72.7% ±2.2
	0.20 $n=4$	+0.0844 ±0.0054	+0.1286 ±0.0030	65.6% ±2.7	73.9% ±3.9
	0.10 $n=4$	+0.0726 ±0.0033	+0.1237 ±0.0036	58.7% ±1.5	71.0% ±2.0
	0.05 $n=4$	+0.0587 ±0.0074	+0.1048 ±0.0061	55.9% ±4.6	73.7% ±3.6

Mean ± sample SD over the data splits scored for that setting; n is on each row. Every entry is a difference from that split's own pooled generalist, because absolute NLL is not comparable across splits. Shading is min-max within a column only. Capture is (router – pooled)/(oracle – pooled): the share of the available headroom the router actually took. It peaks at $\sigma/\hat{\sigma}_0^* = 0.20$ and collapses at 0.70, where the specialists stop separating.