

Mean L2 per pair, split 0. These are distances, not scores: neither column has a good direction.

	trained – approx	trained – empirical	approx – empirical	trained radius
reach 0.05 ($K=7$)	0.395	0.075	0.421	0.564
reach 0.50 ($K=7$)	0.253	0.038	0.275	0.518
reach 0.70 ($K=7$)	0.270	0.123	0.351	0.357
2 pairs ($K=2$)	0.290	0.180	0.442	0.253
5 pairs ($K=5$)	0.240	0.149	0.341	0.463

Pairs are Hungarian-matched before differencing, because the two solves index pairs independently. The space has diagonal $\sqrt{7}$.

Every run was initialised from the approximate solve; the empirical- B solve is offline comparison only. Shaded columns are min-max within the column.