




Tier 1 — per-round SFT training loss (legend shows mean optimiser steps per round)



Tier 3a — cross-NLL at round 9



Tier 3b — fit on own vs others' batches



Tier 3c — specialisation margin (positive = specialised)

